

Exploring propr

David Lovell

December 5, 2016

What's in the Cane Toad data?

`caneToad.counts` contains 20 observations (toads) of the counts of 57580 variables (mRNAs). Each observation is associated with a location:

```
table(caneToad.groups)
```

```
## caneToad.groups
## QLD  WA
##  10  10
```

Zero replacement

Thom kept the RNAs that had at least 10 counts in at least 10 toads

```
keep <- apply(caneToad.counts, 2, function(x) sum(x >= 10) >= 10)
# This line would keep only the RNAs that have at least 10 counts in _all_ toads
# keep <- apply(caneToad.counts, 2, function(x) min(x) >= 10)
counts <- caneToad.counts[,keep]
```

This means we keep:

```
addmargins(table(keep))
```

```
## keep
## FALSE  TRUE  Sum
## 31806 25774 57580
```

Still, of the 515480 counts in this subset of the cane toad data, there are still 1427 zeros. The routines in `propr` automatically replace zero counts with 1. However, no such imputation happens for the scripts that I have developed. So, here I make this replacement explicitly

```
counts[counts == 0] <- 1
```

Transform and rearrange

Take the centred log ratio of counts, put that in a dataframe with the observation ID and location:

```
counts.clr <- clr(counts, check=TRUE)

counts.wide <- data.frame(
  ID=factor(1:nrow(counts)),
  Location=factor(caneToad.groups),
  counts.clr)
counts.long <- gather(counts.wide, RNA, count, -c(ID, Location))
```

Reorder the RNA factor


Set the levels of the RNA factor in descending order of mean count; this is for plotting RNA counts in a more comprehensible way

```
counts.long$RNA <- reorder(counts.long$RNA, -counts.long$count, mean)
#counts.long$RNA <- reorder(counts.long$RNA,
# ifelse(counts.long$Location=="QLD", -counts.long$count, 0), mean)
```

Look at the data

If we plot the sorted RNA counts in chunks we can see everything in a few pages

```
plot4k <- function(n=0){
  ggplot(
    subset(counts.long, n*4000 < as.numeric(RNA) & as.numeric(RNA) <= (n+1)*4000),
    aes(x=1 + (as.numeric(RNA) -1) %% 500, y=count, group=ID)) +
    geom_line(aes(colour=Location), alpha=0.25) +
    facet_grid(((as.numeric(RNA)-1) %% 500 * 500) ~., scales="free_y") +
    labs(title=sprintf("RNAs %d to %d", n*4000 + 1, (n+1)*4000),
         x="RNA sorted in descending order of mean clr across observations",
         y="clr(count)")
}
```



Check that Quinn and Lovell's implementations of propr agree

We sourced Lovell's implementation from `./propr-functions.R`. Now we are going to check to ensure that results from Quinn's and Lovell's implementations agree.

Let's just work with a few components

```
counts.sub <- counts[,1:4]
```

Lovell's implementation uses `phiDF()` to return a list of β , p , r^2 , $\text{var}(\log(x_i, x_j))$, ϕ and now ϕ_s . These results are returned as the lower triangles of the matrices to save space:

```
counts.phiDF <- phiDF(counts.sub)
```

Quinn's implementation uses `phis()`, `phit()` and `perb()` to return ϕ_s , ϕ and ρ_p respectively:

```
counts.phis <- phis(counts.sub)
counts.phit <- phit(counts.sub)
counts.perb <- perb(counts.sub)
```

The next three comparisons check for equality between ϕ_s , ϕ and ρ_p as calculated by Lovell's and Quinn's implementations. Note that we have to compare the lower triangle of the matrices:

```
all.equal(
  counts.phis@matrix[lower.tri(counts.phis@matrix)],
  counts.phiDF$phisym)
```

```
## [1] TRUE
```

```
all.equal(
  counts.phit@matrix[lower.tri(counts.phit@matrix)],
  counts.phiDF$phi)
```

```
## [1] TRUE
```

```
all.equal(
  counts.perb@matrix[lower.tri(counts.perb@matrix)],
  (1 - counts.phiDF$phisym)/(1 + counts.phiDF$phisym))
```

```
## [1] TRUE
```

While it's good that these results agree, let's check further by calculating ϕ_s , and ϕ directly.

```
counts.sub.clr <- clr(counts.sub)
CovXY <- var(counts.sub.clr)
VarX.Plus.VarY <- outer(diag(CovXY), diag(CovXY), "+")
D <- ncol(counts.sub.clr)
VarX <- matrix(rep(diag(CovXY), D), ncol = D, byrow = TRUE)
```

```
all.equal(
  (VarX.Plus.VarY - 2 * CovXY)/(VarX.Plus.VarY + 2 * CovXY),
  phis(counts.sub)@matrix,
  check.attributes=FALSE
)
```

```
## [1] TRUE
```

```
all.equal(
  ((VarX.Plus.VarY - 2 * CovXY)/VarX)[lower.tri(VarX)],
```

```
phit(counts.sub)@matrix[lower.tri(VarX)]  
)
```

```
## [1] TRUE
```

Continue with the cane toad analysis

Calculate ρ_p and ϕ_s

Here we calculate both $\rho_p(\text{clr}(x_i), \text{clr}(x_j))$ and $\phi_s(\text{clr}(x_i), \text{clr}(x_j))$ for the cane toad counts. These

```
rho.p <- perb(caneToad.counts, select = keep)
```

```
## Alert: Replacing 0s in "count matrix" with 1.
```

```
phi.s <- phis(caneToad.counts, select = keep)
```

```
## Alert: Replacing 0s in "count matrix" with 1.
```

Since these should be related monotonically by

$$\phi_s = \frac{1 - \rho_p}{1 + \rho_p}$$

we should end up with the same pairs when we apply the following thresholds:

```
best.rho.p <- rho.p[">", 0.995]
best.phi.s <- phi.s["<", (1 - 0.995)/(1 + 0.995)]
top.rho.p <- simplify(best.rho.p)
top.phi.s <- simplify(best.phi.s)
```

...and indeed we do:

```
all.equal(
  top.rho.p@pairs, top.phi.s@pairs
)
```

```
## [1] TRUE
```


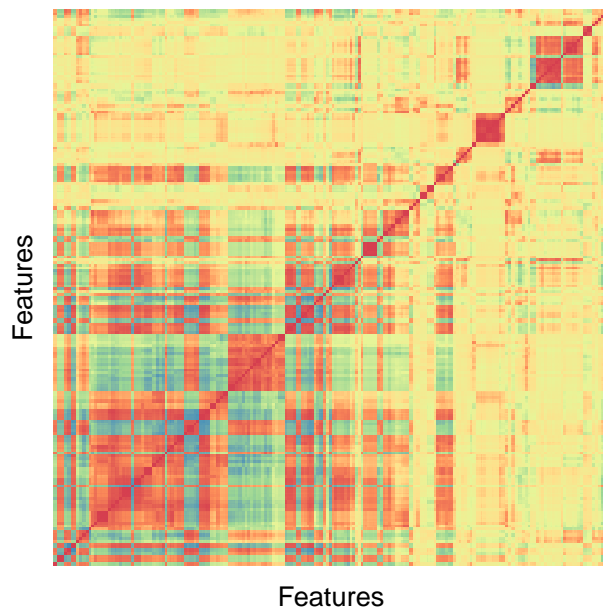

Visualise ρ_p and ϕ_s

Clustered heatmaps

Here are the clustered heatmaps for these two different distances. Note that the dendrogram of the ρ_p matrix shows less distinct clusters than the dendrogram of the ϕ_s matrix. I think this is because the values of ρ_p are constrained within the hypercube $[0, 1]^D$ while the values of ϕ_s can range over $[0, \infty)^D$.

```
dendrogram(top.rho.p)
```




```
## Alert: Generating plot using indexed feature pairs.
```



```
## 'dendrogram' with 2 branches and 403 members total, at height 0.9999977
```

```
dendrogram(top.phi.s)
```

```
## Alert: Generating plot using indexed feature pairs.
```





```
## 'dendrogram' with 2 branches and 403 members total, at height 62.5556
```


PCA projections

Here are the projections of the ρ_p and ϕ_s matrices onto their first two principal components. I'm not sure what these tell us about the data though. First comes the projection performed by `propr::mds()`

```
mds(top.rho.p, group=caneToad.groups)
```

```
## Alert: This function ignores index.
```

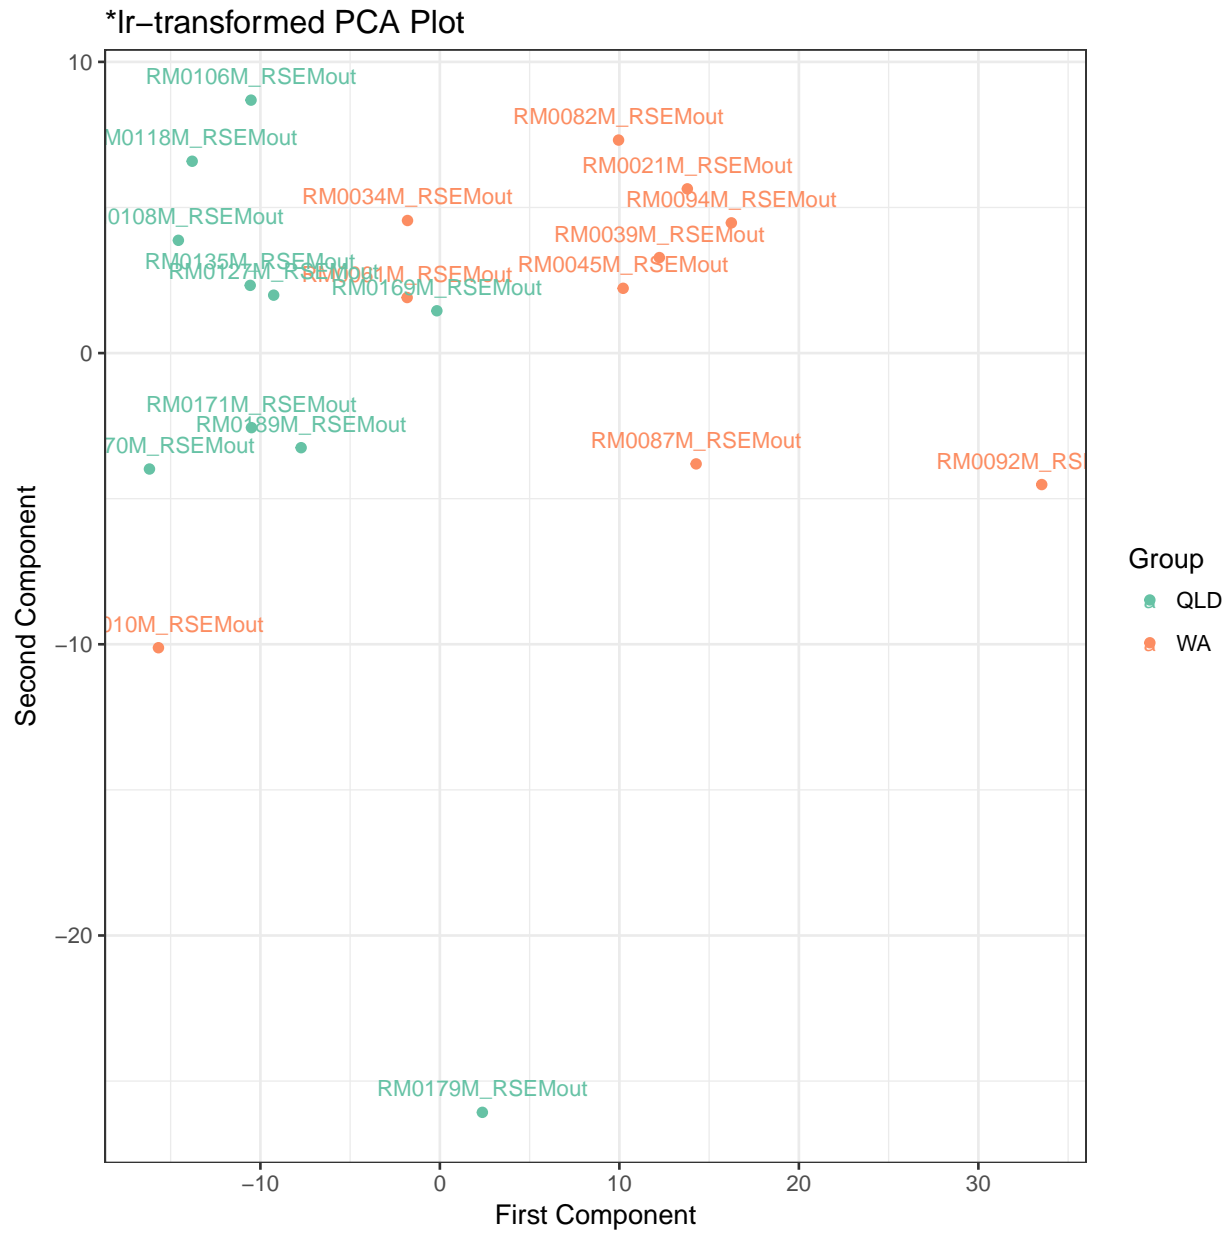


We have to butcher `propr::mds()` a bit to create a function to do the PCA projection of the ϕ_s matrix:

```
pca <- function(object, group){
  if (missing(group)) {
    group <- rep("None", nrow(object@logratio))
  }
  df <- data.frame(
    ID = rownames(object@logratio),
    Group = as.character(group),
    prcomp(object@logratio)$x[, c(1, 2)])
  g <- ggplot2::ggplot(ggplot2::aes_string(ID = "ID"), data = df) +
    ggplot2::geom_point(ggplot2::aes_string(x = "PC1", y = "PC2",
                                             colour = "Group")) +

    ggplot2::theme_bw() +
    ggplot2::xlab("First Component") +
    ggplot2::ylab("Second Component") +
    ggplot2::scale_colour_brewer(palette = "Set2", name = "Group") +
    ggplot2::ggtitle("*lr-transformed PCA Plot") +
    ggplot2::geom_text(
      ggplot2::aes_string(x = "PC1", y = "PC2", label = "ID", colour = "Group"),
      data = df, size = 3, vjust = -1)
  return(g)
}


pca(top.phi.s, group=caneToad.groups)
```



Multidimensional scaling


Here we use `cmdscale()` to perform classical multidimensional scaling (MDS) (also known as principal coordinates analysis (Gower, 1966)) of the ϕ_s matrix:

```
fit <- cmdscale(top.phi.s@matrix, eig=TRUE, k=2)
fit.df <- data.frame(x=fit$points[,1], y=fit$points[,2], mRNA=colnames(top.phi.s@counts))
ggplot(fit.df, aes(x=x, y=y)) +
  geom_point() +
  coord_equal() +
  labs(title="Classical multidimensional scaling of phi_s", caption="caption")
```



```
fit <- cmdscale(1-abs(top.rho.p@matrix), eig=TRUE, k=2)
fit.df <- data.frame(x=fit$points[,1], y=fit$points[,2], mRNA=colnames(top.rho.p@counts))
ggplot(fit.df, aes(x=x, y=y)) +
  geom_point() +
  coord_equal() +
  labs(title="Classical multidimensional scaling of rho_p", caption="caption")
```

Classical multidimensional scaling of rho_p




Graph layout

I'm going to use the output of `phiDF()` to illustrate how I would do a graph layout of the strongly proportional mRNAs.

```
counts.phiDF <- phiDF(top.phi.s@counts)
```

This next graph connects all the mRNA pairs that are strongly proportional with an edge.

```
require(igraph)
g <- graph.data.frame(
  subset(counts.phiDF, phisym < (1 - 0.995)/(1 + 0.995)),
  directed=FALSE)
plot(
  g,
  layout=layout_fruchterman_reingold(g, weight=1/E(g)$phisym),
  vertex.size=1,
  vertex.color="black",
  vertex.label=NA
)
```



We can see that there are duos, trios, quartets and a few larger sets of mRNAs that behave proportionally across the 20 different individuals.